

Lab 6: Introduction to Docker Networking

The goal of this lab is to provide you with a concise overview of how container networking works, how it differs from networking at the level of the Docker host, and how containers can leverage Docker networking to provide direct network connectivity to other containerized services. By the end of this lab, you will know how to deploy containers using networking configurations such as `bridge`, `overlay`, and `host`. You will learn the benefits of different networking drivers and under which circumstances you should choose certain network drivers. Finally, we will look at containerized networking between hosts deployed in a Docker swarm cluster.

Exercise 6.01: Hands-On with Docker Networking

In this exercise, you will run Docker containers and use basic networking to run two simple web servers (Apache2 and NGINX) that will expose ports in a few different basic networking scenarios. You will then access the exposed ports of the container to learn more about how Docker networking works at the most basic level. Launching containers and exposing the service ports to make them available is one of the most common networking scenarios when first starting with containerized infrastructure:

1. List the networks that are currently configured in your Docker environment using the `docker network ls` command:

```
docker network ls
```

The output displayed will show all the configured Docker networks available on your system. It should resemble the following:

NETWORK ID	NAME	DRIVER	SCOPE
0774bdf6228d	bridge	bridge	local
f52b4a5440ad	host	host	local
9bed60b88784	none	null	local

2. When creating a container using Docker without specifying a network or networking driver, Docker will create the container using a `bridge` network. This network exists behind a `bridge` network interface configured in your host OS. Use `ipconfig` to see which interface the Docker bridge is configured as. It is generally called `docker0`:

```
ipconfig
```

The output of this command will list all the network interfaces available in your environment, as shown in the following figure:

```

docker0: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500
    inet 172.17.0.1 netmask 255.255.0.0 broadcast 172.17.255.255
    ether 02:42:c5:b0:7a:34 txqueuelen 0 (Ethernet)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

enpls0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 192.168.122.185 netmask 255.255.255.0 broadcast 192.168.122.255
    inet6 fe80::fcc1:6453:cb2f:37da prefixlen 64 scopeid 0x20<link>
    ether 52:54:00:0a:6a:c1 txqueuelen 1000 (Ethernet)
    RX packets 11735 bytes 140578983 (140.5 MB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 5826 bytes 428029 (428.0 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 221 bytes 17799 (17.7 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 221 bytes 17799 (17.7 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

```

It can be observed in the preceding figure that the Docker bridge interface is called `docker0` and has an IP address of `172.17.0.1`.

- Use the `docker run` command to create a simple NGINX web server container, using the `latest` image tag. Set the container to start in the background using the `-d` flag and give it a human-readable name of `webserver1` using the `--name` flag:

```
docker run -d --name webserver1 nginx:latest
```

If the command is successful, no output will be returned in the terminal session.

- Execute the `docker ps` command to check whether the container is up and running:

```
docker ps
```

As you can see, the `webserver1` container is up and running as expected:

CONTAINER ID	IMAGE	COMMAND	CREATED
STATUS		PORTS	NAMES
0774bdf6228d	nginx:latest	"nginx -g 'daemon of..."	4 seconds ago
Up 3 seconds		80/tcp	webserver1

- Execute the `docker inspect` command to check what networking configuration this container has by default:

```
docker inspect webserver1
```

Docker will return the verbose details about the running container in JSON format. For this exercise, focus on the `NetworkSettings` block. Pay special attention to the `Gateway`, `IPAddress`, `Ports`, and `NetworkID` parameters underneath the `networks` sub-block:

```

"NetworkSettings": {
  "Bridge": "",
  "SandboxID": "d8f43327a8b73a502a1fd0b3e33ac0136d324ea75b4ec73cf453c033ead7c9ba",
  "HairpinMode": false,
  "LinkLocalIPv6Address": "",
  "LinkLocalIPv6PrefixLen": 0,
  "Ports": {
    "80/tcp": null
  },
  "SandboxKey": "/var/run/docker/netns/d8f43327a8b7",
  "SecondaryIPAddresses": null,
  "SecondaryIPv6Addresses": null,
  "EndpointID": "cf88feaa92e4213034102159a4876b99d36ca707b9497aba8468678248bc4a40",
  "Gateway": "172.17.0.1",
  "GlobalIPv6Address": "",
  "GlobalIPv6PrefixLen": 0,
  "IPAddress": "172.17.0.2",
  "IPPrefixLen": 16,
  "IPv6Gateway": "",
  "MacAddress": "02:42:ac:11:00:02",
  "Networks": {
    "bridge": {
      "IPAMConfig": null,
      "Links": null,
      "Aliases": null,
      "NetworkID": "0774bdf6228d30a4758bc972566ac05b80627096d4694e96a3010dd18141aa",
      "EndpointID": "cf88feaa92e4213034102159a4876b99d36ca707b9497aba8468678248bc4a40",
      "Gateway": "172.17.0.1",
      "IPAddress": "172.17.0.2",
      "IPPrefixLen": 16,
      "IPv6Gateway": "",
      "GlobalIPv6Address": "",
      "GlobalIPv6PrefixLen": 0,
      "MacAddress": "02:42:ac:11:00:02",
      "DriverOpts": null
    }
  }
}
}
}

```

8. To expose your service on the network to other users or servers, use the `-p` flag in the `docker run` command. This will allow you to map a port on the host to an exposed port on the container. This is similar to port forwarding on a router or other network device. To expose a container by the port to the outside world, use the `docker run` command followed by the `-d` flag to start the container in the background. The `-p` flag will enable you to specify a port on the host, separated by a colon and the port on the container that you wish to expose. Also, give this container a unique name, `webserver2`:

```
docker run -d -p 8080:80 --name webserver2 nginx:latest
```

Upon successful container startup, your shell will not return anything. However, certain versions of Docker may show the full container ID.

9. Run the `docker ps` command to check whether you have two NGINX containers up and running:

```
docker ps
```

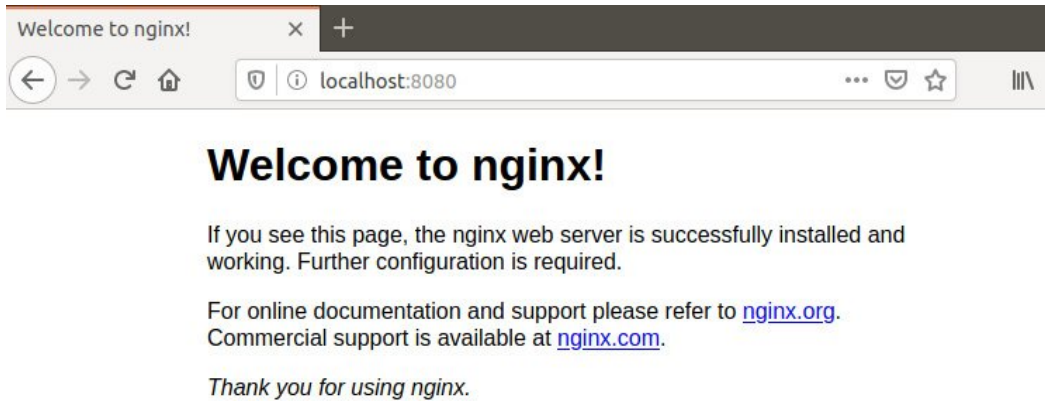
The two running containers, `webserver1` and `webserver2`, will be displayed:

CONTAINER ID	IMAGE	COMMAND	CREATED
STATUS	PORTS	NAMES	
b945fa75b59a	nginx:latest	"nginx -g 'daemon of...'"	1 minute ago
Up About a minute	0.0.0.0:8080->80/tcp	webserver2	

```
3267bf4322ed nginx:latest "nginx -g 'daemon of..." 2 minutes ago
Up 2 minutes          80/tcp          webserver1
```

In the `PORTS` column, you will see that Docker is now forwarding port `80` on the `webserver` container to port `8080` on the host machine. That is deduced from the `0.0.0.0:8080->80/tcp` part of the output.

10. In your web browser, navigate to `http://localhost:8080` to see the running container instance you just spawned:



11. Now, you have two NGINX instances running in the same Docker environment with slightly different networking configurations. The `webserver1` instance is running solely on the Docker network without any ports exposed. Inspect the configuration of the `webserver2` instance using the `docker inspect` command followed by the container name or ID:

```
docker inspect webserver2
```

The `NetworkSettings` section at the bottom of the JSON output will resemble the following. Pay close attention to the parameters (`Gateway` , `IPAddress` , `Ports` , and `NetworkID`) underneath the `networks` sub-block:

```

"NetworkSettings": {
  "Bridge": "",
  "SandboxID": "e271e27f8ea8855a574ed4481dfb0bad9de1da8924ee36cbce17ebabe03773ed",
  "HairpinMode": false,
  "LinkLocalIPv6Address": "",
  "LinkLocalIPv6PrefixLen": 0,
  "Ports": {
    "80/tcp": [
      {
        "HostIp": "0.0.0.0",
        "HostPort": "8080"
      }
    ]
  },
  "SandboxKey": "/var/run/docker/netns/e271e27f8ea8",
  "SecondaryIPAddresses": null,
  "SecondaryIPv6Addresses": null,
  "EndpointID": "55dd352f2dd5c1f096edb2c4c65fb7906c8201a5e89b32668904b08f51a2c4b4",
  "Gateway": "172.17.0.1",
  "GlobalIPv6Address": "",
  "GlobalIPv6PrefixLen": 0,
  "IPAddress": "172.17.0.3",
  "IPPrefixLen": 16,
  "IPv6Gateway": "",
  "MacAddress": "02:42:ac:11:00:03",
  "Networks": {
    "bridge": {
      "IPAMConfig": null,
      "Links": null,
      "Aliases": null,
      "NetworkID": "b1d6e724b189ee7fb6bd66d0152d9cd1f50787b82e423f4ff1da17472861e420",
      "EndpointID": "55dd352f2dd5c1f096edb2c4c65fb7906c8201a5e89b32668904b08f51a2c4b4",
      "Gateway": "172.17.0.1",
      "IPAddress": "172.17.0.3",
      "IPPrefixLen": 16,
      "IPv6Gateway": "",
      "GlobalIPv6Address": "",
      "GlobalIPv6PrefixLen": 0,
      "MacAddress": "02:42:ac:11:00:03",
      "DriverOpts": null
    }
  }
}

```

As the `docker inspect` output displays, the `webserver2` container has an IP address of `172.17.0.3`, whereas your `webserver1` container has an IP address of `172.17.0.1`. The IP addresses in your local environment may be slightly different depending on how Docker assigns the IP addresses to the containers. Both the containers live on the same Docker network (`bridge`) and have the same default gateway, which is the `docker0` `bridge` interface on the host machine.

- Since both of these containers live on the same subnet, you can test communication between the containers within the Docker `bridge` network. Run the `docker exec` command to gain access to a shell on the `webserver1` container:

```
docker exec -it webserver1 /bin/bash
```

The prompt should noticeably change to a root prompt, indicating you are now in a Bash shell on the `webserver1` container:

```
root@3267bf4322ed:/#
```

- At the root shell prompt, use the `apt` package manager to install the `ping` utility in this container:

```
root@3267bf4322ed:/# apt-get update && apt-get install -y inetutils-ping
```

The aptitude package manager will then install the `ping` utility in the `webserver1` container. Please note that the `apt` package manager will install `ping` as well as other dependencies that are required to run the `ping` command:

```
root@3267bf4322ed:/# apt-get update && apt-get install -y inetutils-ping
Hit:1 http://deb.debian.org/debian buster InRelease
Hit:2 http://security.debian.org/debian-security buster/updates InRelease
Hit:3 http://deb.debian.org/debian buster-updates InRelease
Reading package lists... Done
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following additional packages will be installed:
  libidn11 netbase
The following NEW packages will be installed:
  inetutils-ping libidn11 netbase
0 upgraded, 3 newly installed, 0 to remove and 0 not upgraded.
Need to get 362 kB of archives.
After this operation, 718 kB of additional disk space will be used.
Get:1 http://deb.debian.org/debian buster/main amd64 netbase all 5.6 [19.4 kB]
Get:2 http://deb.debian.org/debian buster/main amd64 libidn11 amd64 1.33-2.2 [116 kB]
Get:3 http://deb.debian.org/debian buster/main amd64 inetutils-ping amd64 2:1.9.4-7 [226 kB]
Fetched 362 kB in 0s (3133 kB/s)
debconf: delaying package configuration, since apt-utils is not installed
Selecting previously unselected package netbase.
(Reading database ... 7203 files and directories currently installed.)
Preparing to unpack .../archives/netbase_5.6_all.deb ...
Unpacking netbase (5.6) ...
Selecting previously unselected package libidn11:amd64.
Preparing to unpack .../libidn11_1.33-2.2_amd64.deb ...
Unpacking libidn11:amd64 (1.33-2.2) ...
Selecting previously unselected package inetutils-ping.
Preparing to unpack .../inetutils-ping_2%3a1.9.4-7_amd64.deb ...
Unpacking inetutils-ping (2:1.9.4-7) ...
Setting up libidn11:amd64 (1.33-2.2) ...
Setting up netbase (5.6) ...
Setting up inetutils-ping (2:1.9.4-7) ...
Processing triggers for libc-bin (2.28-10) ...
```

14. Once the `ping` utility has successfully installed, use it to ping the IP address of the other container:

```
root@3267bf4322ed:/# ping <ADD_WEBSERVER2_IP_ADDRESS_HERE>
```

```
Administrator: Command Prompt - docker exec -it webserver1 /bin/bash
(Reading database ... 7815 files and directories currently installed.)
Preparing to unpack .../archives/netbase_6.3_all.deb ...
Unpacking netbase (6.3) ...
Selecting previously unselected package inetutils-ping.
Preparing to unpack .../inetutils-ping_2%3a2.0-1_amd64.deb ...
Unpacking inetutils-ping (2:2.0-1) ...
Setting up netbase (6.3) ...
Setting up inetutils-ping (2:2.0-1) ...
root@8f2a225057a4:/# ping 172.17.0.5
PING 172.17.0.5 (172.17.0.5): 56 data bytes
64 bytes from 172.17.0.5: icmp_seq=0 ttl=64 time=0.169 ms
64 bytes from 172.17.0.5: icmp_seq=1 ttl=64 time=0.124 ms
64 bytes from 172.17.0.5: icmp_seq=2 ttl=64 time=0.323 ms
64 bytes from 172.17.0.5: icmp_seq=3 ttl=64 time=0.096 ms
64 bytes from 172.17.0.5: icmp_seq=4 ttl=64 time=0.110 ms
64 bytes from 172.17.0.5: icmp_seq=5 ttl=64 time=0.141 ms
^C--- 172.17.0.5 ping statistics ---
6 packets transmitted, 6 packets received, 0% packet loss
round-trip min/avg/max/stddev = 0.096/0.161/0.323/0.076 ms
```


The output should display ICMP response packets, indicating that the containers can successfully ping each other through the Docker `bridge` network:

```
PING 172.17.0.1 (172.17.0.3): 56 data bytes
64 bytes from 172.17.0.3: icmp_seq=0 ttl=64 time=0.221 ms
64 bytes from 172.17.0.3: icmp_seq=1 ttl=64 time=0.207 ms
```

15. You can also access the NGINX default web interface using the `curl` command. Install `curl` using the `apt` package manager:

```
root@3267bf4322ed:/# apt-get install -y curl
```

The following output should display, indicating that the `curl` utility and all required dependencies are being installed:

```
Building dependency tree
Reading state information... Done
Need to get 3471 kB of archives.
After this operation, 7725 kB of additional disk space will be used.
Get:1 http://deb.debian.org/debian buster/main amd64 krb5-locales all 1.17-3 [95.4 kB]
Get:2 http://security.debian.org/debian-security buster/updates/main amd64 libldap-common all 2.4.47+dfsg-3+deb10u2 [89.7 kB]
Get:3 http://deb.debian.org/debian buster/main amd64 libsasl2-modules-db amd64 2.1.27+dfsg-1+deb10u1 [69.1 kB]
Get:4 http://deb.debian.org/debian buster/main amd64 libsasl2-2 amd64 2.1.27+dfsg-1+deb10u1 [106 kB]
Get:5 http://deb.debian.org/debian buster/main amd64 ca-certificates all 20190110 [157 kB]
Get:6 http://deb.debian.org/debian buster/main amd64 libkeyutils1 amd64 1.6-6 [15.0 kB]
Get:7 http://deb.debian.org/debian buster/main amd64 libkrb5support0 amd64 1.17-3 [65.6 kB]
Get:8 http://deb.debian.org/debian buster/main amd64 libk5crypto3 amd64 1.17-3 [121 kB]
Get:9 http://deb.debian.org/debian buster/main amd64 libkrb5-3 amd64 1.17-3 [370 kB]
Setting up libsasl2-modules:amd64 (2.1.27+dfsg-1+deb10u1) ...
Setting up libnghttp2-14:amd64 (1.36.0-2+deb10u1) ...
Setting up krb5-locales (1.17-3) ...
Setting up libldap-common (2.4.47+dfsg-3+deb10u2) ...
Setting up libkrb5support0:amd64 (1.17-3) ...
Setting up libsasl2-modules-db:amd64 (2.1.27+dfsg-1+deb10u1) ...
Setting up librtmp1:amd64 (2.4+20151223.gitfa8646d.1-2) ...
Setting up libk5crypto3:amd64 (1.17-3) ...
Setting up libsasl2-2:amd64 (2.1.27+dfsg-1+deb10u1) ...
Setting up libssh2-1:amd64 (1.8.0-2.1) ...
Setting up libkrb5-3:amd64 (1.17-3) ...
Setting up openssl (1.1.1d-0+deb10u3) ...
Setting up publicsuffix (20190415.1030-1) ...
Setting up libldap-2.4-2:amd64 (2.4.47+dfsg-3+deb10u2) ...
Setting up ca-certificates (20190110) ...
debconf: unable to initialize frontend: Dialog
Updating certificates in /etc/ssl/certs...
128 added, 0 removed; done.
Setting up libgssapi-krb5-2:amd64 (1.17-3) ...
Setting up libcurl4:amd64 (7.64.0-4+deb10u1) ...
Setting up curl (7.64.0-4+deb10u1) ...
Processing triggers for libc-bin (2.28-10) ...
Processing triggers for ca-certificates (20190110) ...
Updating certificates in /etc/ssl/certs...
0 added, 0 removed; done.
Running hooks in /etc/ca-certificates/update.d...
done
```

16. After installing `curl`, use it to curl the IP address of `webserver2`:

```
root@3267bf4322ed:/# curl <ADD_WEBSEVER2_IP_ADDRESS_HERE>
```

You should see the `Welcome to nginx!` page displayed in HTML format, indicating that you were able to successfully contact the IP address of the `webserver2` container through the Docker `bridge` network:

```
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
```

```

    body {
        width: 35em;
        margin: 0 auto;
        font-family: Tahoma, Verdana, Arial, sans-serif;
    }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully
installed and working. Further configuration is required.</p>
<p>For online documentation and support please refer to
<a href="http://nginx.org/">nginx.org</a>.<br/>
Commercial support is available at
<a href="http://nginx.com/">nginx.com</a>.</p>
<p><em>Thank you for using nginx.</em></p>
</body>
</html>

```

Since you are using `curl` to navigate to the NGINX welcome page, it will render on your terminal display in raw HTML format.

During this exercise, we saw how containers can use Docker networks to talk to other containers directly. We used the `webserver1` container to call the IP address of the `webserver2` container and display the output of the web page the container was hosting.

In this exercise, we were also able to demonstrate network connectivity between container instances using the native Docker `bridge` network. However, when we deploy containers at scale, there is no easy way to know which IP address in the Docker network belongs to which container.

In the next section, we will look at native Docker DNS and learn how to use human-readable DNS names to reliably send network traffic to other container instances.

Exercise 6.02: Working with Docker DNS

In the following exercise, you will learn about name resolution between Docker containers running on the same network. You will first enable simple name resolution using the legacy link method. You will contrast this approach by using the newer and more reliable native Docker DNS service:

1. First, create two Alpine Linux containers on the default Docker `bridge` network that will communicate with each other using the `--link` flag. Alpine is a very good base image for this exercise because it contains the `ping` utility by default. This will enable you to quickly test the connectivity between containers in the various scenarios. To get started, create a container called `containerlink1` to indicate that you have created this container using the legacy link method:

```
docker run -itd --name containerlink1 alpine:latest
```

This will start a container in the default Docker network called `containerlink1`.

2. Start another container in the default Docker bridge network, called `containerlink2`, which will create a link to `containerlink1` to enable rudimentary DNS:

```
docker run -itd --name containerlink2 --link containerlink1 alpine:latest
```


This will start a container in the default Docker network called `containerlink2`.

3. Run the `docker exec` command to access a shell inside the `containerlink2` container. This will allow you to investigate how the link functionality is working. Since this container is running Alpine Linux, you do not have access to the Bash shell by default. Instead, access it using an `sh` shell:

```
docker exec -it containerlink2 /bin/sh
```

This should drop you into a root `sh` shell in the `containerlink2` container.

4. From the shell of the `containerlink2` container, ping `containerlink1`:

```
/ # ping containerlink1
```

You will get a reply to the `ping` request:

```
PING container1 (172.17.0.X): 56 data bytes
64 bytes from 172.17.0.X: seq=0 ttl=64 time=0.307 ms
64 bytes from 172.17.0.X: seq=1 ttl=64 time=0.162 ms
64 bytes from 172.17.0.X: seq=2 ttl=64 time=0.177 ms
```

5. Use the `cat` utility to have a look at the `/etc/hosts` file of the `containerlink2` container. The `hosts` file is a list of routable names to IP addresses that Docker can maintain and override:

```
/ # cat /etc/hosts
```

The output of the `hosts` file should display and resemble the following:

```
127.0.0.1    localhost
::1         localhost ip6-localhost ip6-loopback
fe00::0     ip6-localnet
ff00::0     ip6-mcastprefix
ff02::1     ip6-allnodes
ff02::2     ip6-allrouters
172.17.0.2   containerlink1 032f038abfba
172.17.0.3   9b62c4a57ce3
```

From the output of the `hosts` file of the `containerlink2` container, observe that Docker is adding an entry for the `containerlink1` container name as well as its container ID. This enables the `containerlink2` container to know the name, and the container ID is mapped to the IP address. Typing the `exit` command will terminate the `sh` shell session and bring you back to your environment's main terminal.

6. Run `docker exec` to access an `sh` shell inside the `containerlink1` container:

```
docker exec -it containerlink1 /bin/sh
```

This should drop you into the shell of the `containerlink1` container.

7. Ping the `containerlink2` container using the `ping` utility:

```
/ # ping containerlink2
```

You should see the following output:

```
ping: bad address 'containerlink2'
```

It is not possible to ping the `containerlink2` container since linking containers only works unidirectionally. The `containerlink1` container has no idea that the `containerlink2` container exists since no `hosts` file entry has been created in the `containerlink1` container instance.

8. Due to the limitations using the legacy link method, Docker also supports native DNS using user-created Docker networks. To leverage this functionality, create a Docker network called `dnsnet` and deploy two Alpine containers within that network. First, use the `docker network create` command to create a new Docker network using a `192.168.56.0/24` subnet and using the IP address `192.168.54.1` as the default gateway:

```
docker network create dnsnet --subnet 192.168.54.0/24 --gateway 192.168.54.1
```

Note

Simply using the `docker network create dnsnet` command will create a network with a Docker-allocated subnet and gateway. This exercise demonstrates how to specify the subnet and gateway for your Docker network. It should also be noted that if your computer is attached to a subnet in the `192.168.54.0/24` subnet or a subnet that overlaps that space, it may cause network connectivity issues. Please use a different subnet for this exercise.

9. Use the `docker network ls` command to list the Docker networks available in this environment:

```
docker network ls
```

The list of Docker networks should be returned, including the `dnsnet` network you just created:

NETWORK ID	NAME	DRIVER	SCOPE
ec5b91e88a6f	bridge	bridge	local
c804e768413d	dnsnet	bridge	local
f52b4a5440ad	host	host	local
9bed60b88784	none	null	local

10. Run the `docker network inspect` command to view the configuration for this network:

```
docker network inspect dnsnet
```

The details of the `dnsnet` network should be displayed. Pay close attention to the `Subnet` and `Gateway` parameters. These are the same parameters that you used to create a Docker network in *Step 8*:

```
[
  {
    "Name": "dnsnet",
    "Id": "c804e768413d14686d570451c8ef34b18a8bb975a1e4994a29cd9964ae8cd6ca",
    "Created": "2020-05-12T19:31:39.300406513-04:00",
    "Scope": "local",
    "Driver": "bridge",
    "EnableIPv6": false,
    "IPAM": {
      "Driver": "default",
      "Options": {},
      "Config": [
        {
          "Subnet": "192.168.54.0/24",
          "Gateway": "192.168.54.1"
        }
      ]
    },
    "Internal": false,
    "Attachable": false,
    "Ingress": false,
    "ConfigFrom": {
      "Network": ""
    },
    "ConfigOnly": false,
    "Containers": {},
    "Options": {},
    "Labels": {}
  }
]
```

11. Now that a new Docker network has been created, use the `docker run` command to start a new container (`alpinedns1`) within this network. Use the `docker run` command with the `--network` flag to specify the `dnsnet` network that was just created, and the `--network-alias` flag to give your container a custom DNS name:

```
docker run -itd --network dnsnet --network-alias alpinedns1 --name alpinedns1
alpine:latest
```

Upon successful execution of the command, the full container ID should be displayed before returning to a normal terminal prompt.

12. Start a second container (`alpinedns2`) using the same `--network` and `--network-alias` settings:

```
docker run -itd --network dnsnet --network-alias alpinedns2 --name alpinedns2
alpine:latest
```

Note

It is important to understand the difference between the `--network-alias` flag and the `--name` flag. The `--name` flag is used to give the container a human-readable name within the Docker API. This makes it easy to start, stop, restart, and manage containers by name. The `--network-alias` flag, however, is used to create a custom DNS entry for the container.

13. Use the `docker ps` command to verify that the containers are running as expected:

```
docker ps
```

The output will display the running container instances.

15. Use the `docker inspect` command to verify that the IP addresses of the container instances are from within the subnet (`192.168.54.0/24`) that was specified:

```
docker inspect alpinedns1
```

The following output is truncated to show the relevant details:

```
"Networks": {
  "dnsnet": {
    "IPAMConfig": null,
    "Links": null,
    "Aliases": [
      "alpinedns1",
      "69ecb9ad45e1"
    ],
    "NetworkID": "c804e768413d14686d570451c8ef34b18a8bb975a1e4994a29cd9964ae8cd6ca",
    "EndpointID": "c7f9b4f3a6aaabd7118eeefcdd8f11db1af6f10d1c90c0b9eac6a098cda071a",
    "Gateway": "192.168.54.2",
    "IPAddress": "192.168.54.1",
    "IPPrefixLen": 24,
    "IPv6Gateway": "",
    "GlobalIPv6Address": "",
    "GlobalIPv6PrefixLen": 0,
    "MacAddress": "02:42:c0:a8:36:03",
    "DriverOpts": null
  }
}
```

It can be observed from the output that the `alpinedns1` container was deployed with an IP address of `192.168.54.2` , which is a part of the subnet that was defined during the creation of the Docker network.

16. Execute the `docker network inspect` command in a similar fashion for the `alpinedns2` container:

```
docker inspect alpinedns2
```

The output is again truncated to display the relevant networking details:

```
"Networks": {
  "dnsnet": {
    "IPAMConfig": null,
    "Links": null,
    "Aliases": [
      "alpinedns2",
      "69ecb9ad45e1"
    ],
    "NetworkID": "c804e768413d14686d570451c8ef34b18a8bb975a1e4994a29cd9964ae8cd6ca",
    "EndpointID": "c7f9b4f3a6aaabd7118eeefcdd8f11db1af6f10d1c90c0b9eac6a098cda071a",
    "Gateway": "192.168.54.1",
    "IPAddress": "192.168.54.3",
    "IPPrefixLen": 24,
    "IPv6Gateway": "",
    "GlobalIPv6Address": "",
    "GlobalIPv6PrefixLen": 0,
    "MacAddress": "02:42:c0:a8:36:03",
    "DriverOpts": null
  }
}
```

It can be observed in the preceding output that the `alpinedns2` container has an IP address of `192.168.54.3`, which is a different IP address within the `dnsnet` subnet.

17. Run the `docker exec` command to access a shell in the `alpinedns1` container:

```
docker exec -it alpinedns1 /bin/sh
```

This should drop you into a root shell inside of the containers.

18. Once inside the `alpinedns1` container, use the `ping` utility to ping the `alpinedns2` container:

```
/ # ping alpinedns2
```

The `ping` output should display successful network connectivity to the `alpinedns2` container instance:

```
PING alpinedns2 (192.168.54.3): 56 data bytes
64 bytes from 192.168.54.3: seq=0 ttl=64 time=0.278 ms
64 bytes from 192.168.54.3: seq=1 ttl=64 time=0.233 ms
```

19. Use the `exit` command to return to your primary terminal. Use the `docker exec` command to gain access to a shell inside the `alpinedns2` container:

```
docker exec -it alpinedns2 /bin/sh
```

This should drop you to a shell within the `alpinedns2` container.

20. Use the `ping` utility to ping the `alpinedns1` container by name:

```
ping alpinedns1
```

The output should display successful responses from the `alpinedns1` container:

```
PING alpinedns1 (192.168.54.2): 56 data bytes
64 bytes from 192.168.54.2: seq=0 ttl=64 time=0.115 ms
64 bytes from 192.168.54.2: seq=1 ttl=64 time=0.231 ms
```

21. Use the `cat` utility inside any of the `alpinedns` containers to reveal that Docker is using true DNS as opposed to `/etc/hosts` file entries inside the container:

```
# cat /etc/hosts
```

This will reveal the contents of the `/etc/hosts` file inside the respective container:

```
127.0.0.1 localhost
::1 localhost ip6-localhost ip6-loopback
fe00::0 ip6-localnet
ff00::0 ip6-mcastprefix
ff02::1 ip6-allnodes
ff02::2 ip6-allrouters
192.168.54.2 9b57038fb9c8
```

Use the `exit` command to terminate the shell session inside of the `alpinedns2` container.

22. Clean up your environment by stopping all running containers using the `docker stop` command:

```
docker stop containerlink1
docker stop containerlink2
docker stop alpinedns1
docker stop alpinedns2
```

23. Use the `docker system prune -fa` command to clean the remaining stopped containers and networks:

```
docker system prune -fa
```

Successfully executing this command should clean up the `dnsnet` network as well as the container instances and images:

```
Deleted Containers:
69ecb9ad45e16ef158539761edc95fc83b54bd2c0d2ef55abfbala300f141c7c
9b57038fb9c8cf30aaebe6485e9d223041a9db4e94eb1be9392132bdef632067
Deleted Networks:
dnsnet
Deleted Images:
untagged: alpine:latest
untagged: alpine@sha256:9a839e63dad54c3a6d1834e29692c8492d93f90c
59c978c1ed79109ea4fb9a54
deleted: sha256:f70734b6a266dcb5f44c383274821207885b549b75c8e119
404917a61335981a
deleted: sha256:3e207b409db364b595ba862cdc12be96dcdad8e36c59a03b
b3b61c946a5741a
Total reclaimed space: 42.12M
```

Each section of the system prune output will identify and remove Docker resources that are no longer in use. In this case, it will remove the `dnsnet` network since no container instances are currently deployed in this network.

In this exercise, you looked at the benefits of using name resolution to enable communication between the containers over Docker networks. Using name resolution is efficient since applications don't have to worry about the IP addresses of the other running containers. Instead, communication can be initiated by simply calling the other containers by name.

In the next section, we will learn about deploying containers using other types of networking drivers to truly provide maximum flexibility when deploying containerized infrastructure.

Exercise 6.03: Exploring Docker Networks

In this exercise, we will look into the various types of Docker network drivers that are supported in Docker by default, such as `host` and `none`. We will start with the `bridge` network and then look into the `none` and `host` networks:

1. Use the `docker network ls` command to view the networks available in your Docker environment:

```
docker network ls
```


This should list the three basic network types defined previously, displaying the network ID, the name of the Docker network, and the driver associated with the network type:

NETWORK ID	NAME	DRIVER	SCOPE
50de4997649a	bridge	bridge	local
f52b4a5440ad	host	host	local
9bed60b88784	none	null	local

2. View the verbose details of these networks using the `docker network inspect` command, followed by the ID or the name of the network you want to inspect. In this step, you will view the verbose details of the `bridge` network:

```
docker network inspect bridge
```

Docker will display the verbose output of the `bridge` network in JSON format:

```
[
  {
    "Name": "bridge",
    "Id": "50de4997649a36b77ac540f549d9931a64a7009ee2be8489b63a4a220b918b4e",
    "Created": "2020-05-01T20:56:38.522661526-04:00",
    "Scope": "local",
    "Driver": "bridge",
    "EnableIPv6": false,
    "IPAM": {
      "Driver": "default",
      "Options": null,
      "Config": [
        {
          "Subnet": "172.17.0.0/16",
          "Gateway": "172.17.0.1"
        }
      ]
    },
    "Internal": false,
    "Attachable": false,
    "Ingress": false,
    "ConfigFrom": {
      "Network": ""
    },
    "ConfigOnly": false,
    "Containers": {},
    "Options": {
      "com.docker.network.bridge.default_bridge": "true",
      "com.docker.network.bridge.enable_icc": "true",
      "com.docker.network.bridge.enable_ip_masquerade": "true",
      "com.docker.network.bridge.host_binding_ipv4": "0.0.0.0",
      "com.docker.network.bridge.name": "docker0",
      "com.docker.network.driver.mtu": "1500"
    },
    "Labels": {}
  }
]
```

3. View the verbose details of the `host` network using the `docker network inspect` command:

```
docker network inspect host
```

This will display the details of the `host` network in JSON format:

```
[
  {
    "Name": "host",
    "Id": "f52b4a5440ad41e44adb3e14f4892fb57e95479f99d5b67e39739f6bc0a8f13f",
    "Created": "2020-04-11T11:41:59.811741058-04:00",
    "Scope": "local",
    "Driver": "host",
    "EnableIPv6": false,
    "IPAM": {
      "Driver": "default",
      "Options": null,
      "Config": []
    },
    "Internal": false,
    "Attachable": false,
    "Ingress": false,
    "ConfigFrom": {
      "Network": ""
    },
    "ConfigOnly": false,
    "Containers": {},
    "Options": {},
    "Labels": {}
  }
]
```

As you can see, there is not very much configuration present in the `host` network. Since it uses the `host` networking driver, all the container's networking will be shared with the host. Hence, this network configuration does not need to define specific subnets, interfaces, or other metadata, as we have seen in the default `bridge` network from before.

5. Investigate the `none` network next. Use the `docker network inspect` command to view the details of the `none` network:

```
docker network inspect none
```

The details will be displayed in JSON format:

```
[
  {
    "Name": "none",
    "Id": "9bed60b88784312abcae6d6026dec4445abf34964f563e6f1552675a8fd03ab6",
    "Created": "2020-04-11T11:41:59.732455407-04:00",
    "Scope": "local",
    "Driver": "null",
    "EnableIPv6": false,
    "IPAM": {
      "Driver": "default",
      "Options": null,
      "Config": []
    },
    "Internal": false,
    "Attachable": false,
    "Ingress": false,
    "ConfigFrom": {
      "Network": ""
    },
    "ConfigOnly": false,
    "Containers": {},
    "Options": {},
    "Labels": {}
  }
]
```

Similar to the `host` network, the `none` network is mostly empty. Since containers deployed in this network will have no network connectivity by leveraging the `null` driver, there isn't much need for configuration.

- Now create a container in the `none` network to observe its operation. In your terminal or PowerShell session, use the `docker run` command to start an Alpine Linux container in the `none` network using the `--network` flag. Name this container `nonenet` so we know that it is deployed in the `none` network:

```
docker run -itd --network none --name nonenet alpine:latest
```

This will pull and start an Alpine Linux Docker container in the `none` network.

- Use the `docker ps` command to verify whether the container is up and running as expected:

```
docker ps
```

The output should display the `nonenet` container as up and running:

CONTAINER ID	IMAGE	COMMAND	CREATED
STATUS	PORTS	NAMES	
972a80984703	alpine:latest	"/bin/sh"	9 seconds ago
Up 7 seconds		nonenet	

- Execute the `docker inspect` command, along with the container name, `nonenet`, to get a deeper understanding of how this container is configured:

```
docker inspect nonenet
```

The output of `docker inspect` will display the full container configuration in JSON format. A truncated version highlighting the `NetworkSettings` section is provided here. Pay close attention to the `IPAddress` and `Gateway` settings:

```
"NetworkSettings": {
  "Bridge": "",
  "SandboxID": "alf59b07894b3903f11c0c145ed4e650f5cca1003c5fbbac5741552cf87078d3",
  "HairpinMode": false,
  "LinkLocalIPv6Address": "",
  "LinkLocalIPv6PrefixLen": 0,
  "Ports": {},
  "SandboxKey": "/var/run/docker/netns/alf59b07894b",
  "SecondaryIPAddresses": null,
  "SecondaryIPv6Addresses": null,
  "EndpointID": "",
  "Gateway": "",
  "GlobalIPv6Address": "",
  "GlobalIPv6PrefixLen": 0,
  "IPAddress": "",
  "IPPrefixLen": 0,
  "IPv6Gateway": "",
  "MacAddress": "",
  "Networks": {
    "none": {
      "IPAMConfig": null,
      "Links": null,
      "Aliases": null,
      "NetworkID": "9bed60b88784312abcae6d6026dec4445abf34964f563e6f1552675a8fd03ab6",
      "EndpointID": "032b2f248b2f412222129e4597561c7b4b29b11a16a9e65efc906e71437fe14",
      "Gateway": "",
      "IPAddress": "",
      "IPPrefixLen": 0,
      "IPv6Gateway": "",
      "GlobalIPv6Address": "",
      "GlobalIPv6PrefixLen": 0,
      "MacAddress": "",
      "DriverOpts": null
    }
  }
}
```

The `docker inspect` output will reveal that this container does not have an IP address, nor does it have a gateway or any other networking settings.

9. Use the `docker exec` command to access an `sh` shell inside this container:

```
docker exec -it nonenet /bin/sh
```

Upon successful execution of this command, you will be dropped into a root shell in the container instance:

```
/ #
```

10. Execute the `ip a` command to view the network interfaces available in the container:

```
ip a
```

This will display all network interfaces configured in this container:

```
C:\Users\fenago>docker exec -it nonenet /bin/sh
/ # ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
2: tunl0@NONE: <NOARP> mtu 1480 qdisc noop state DOWN qlen 1000
    link/ipip 0.0.0.0 brd 0.0.0.0
3: ip6tnl0@NONE: <NOARP> mtu 1452 qdisc noop state DOWN qlen 1000
    link/tunnel6 00:00:00:00:00:00:00:00:00:00:00:00:00:00:00:00:00:00 brd 00:00:00:00:00:00:00:00:00:00:00:00:00:00:00:00
/ #
```

The only network interface available to this container is its `LOOPBACK` interface. As this container is not configured with an IP address or default gateway, common networking commands will not work.

11. Test the lack of network connectivity using the `ping` utility provided by default in the Alpine Linux Docker image. Try to ping the Google DNS servers located at IP address `8.8.8.8`:

```
/ # ping 8.8.8.8
```

The output of the `ping` command should reveal that it has no network connectivity:

```
PING 8.8.8.8 (8.8.8.8): 56 data bytes
ping: sendto: Network unreachable
```

Use the `exit` command to return to your main terminal session.

12. To get started with running a container in `host` mode. Use the `docker run` command to start an Alpine Linux container in the `host` network. Name it `hostnet1` to tell it apart from the other containers:

```
docker run -itd --network host --name hostnet1 alpine:latest
```

Docker will start this container in the background using the `host` network.

13. Use the `docker inspect` command to look at the network configuration of the `hostnet1` container you just created:

```
docker inspect hostnet1
```

This will reveal the verbose configuration of the running container, including the networking details, in JSON format:

```

"NetworkSettings": {
  "Bridge": "",
  "SandboxID": "67b54051821d44ffd55c7042e51444777fe53ed7b4e5e6a222c485d43321449d",
  "HairpinMode": false,
  "LinkLocalIPv6Address": "",
  "LinkLocalIPv6PrefixLen": 0,
  "Ports": {},
  "SandboxKey": "/var/run/docker/netns/default",
  "SecondaryIPAddresses": null,
  "SecondaryIPv6Addresses": null,
  "EndpointID": "",
  "Gateway": "",
  "GlobalIPv6Address": "",
  "GlobalIPv6PrefixLen": 0,
  "IPAddress": "",
  "IPPrefixLen": 0,
  "IPv6Gateway": "",
  "MacAddress": "",
  "Networks": {
    "host": {
      "IPAMConfig": null,
      "Links": null,
      "Aliases": null,
      "NetworkID": "f52b4a5440ad41e44adb3e14f4892fb57e95479f99d5b67e39739f6bc0a8f13f",
      "EndpointID": "7e7557baccc2cf640f7a8c59e685542f02a442f2531dd0aa7bf02a98735f972b",
      "Gateway": "",
      "IPAddress": "",
      "IPPrefixLen": 0,
      "IPv6Gateway": "",
      "GlobalIPv6Address": "",
      "GlobalIPv6PrefixLen": 0,
      "MacAddress": "",
      "DriverOpts": null
    }
  }
}

```

It should be noted that the output of the `NetworkSettings` block will look a lot like the containers you deployed in the `none` network. In the `host` networking mode, Docker will not assign an IP address or gateway to the container instance since it shares all network interfaces with the host machine directly.

15. Use `docker exec` to access an `sh` shell inside this container, providing the name `hostnet1`:

```
docker exec -it hostnet1 /bin/sh
```

This should drop you into a root shell inside the `hostnet1` container.

16. Inside the `hostnet1` container, execute the `ifconfig` command to list which network interfaces are available to it:

```
/ # ifconfig
```

The full list of network interfaces available inside of this container should be displayed:


```

docker0    Link encap:Ethernet  HWaddr 02:42:80:07:84:96
           inet addr:172.17.0.1  Bcast:172.17.255.255  Mask:255.255.0.0
           UP BROADCAST MULTICAST  MTU:1500  Metric:1
           RX packets:0 errors:0 dropped:0 overruns:0 frame:0
           TX packets:0 errors:0 dropped:0 overruns:0 carrier:0
           collisions:0 txqueuelen:0
           RX bytes:0 (0.0 B)  TX bytes:0 (0.0 B)

enp1s0     Link encap:Ethernet  HWaddr 52:54:00:0A:6A:C1
           inet addr:192.168.122.185  Bcast:192.168.122.255  Mask:255.255.255.0
           inet6 addr: fe80::fcc1:6453:cb2f:37da/64 Scope:Link
           UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
           RX packets:39516 errors:0 dropped:4 overruns:0 frame:0
           TX packets:6201 errors:0 dropped:0 overruns:0 carrier:0
           collisions:0 txqueuelen:1000
           RX bytes:66258728 (63.1 MiB)  TX bytes:542932 (530.2 KiB)

lo         Link encap:Local Loopback
           inet addr:127.0.0.1  Mask:255.0.0.0
           inet6 addr: ::1/128 Scope:Host
           UP LOOPBACK RUNNING  MTU:65536  Metric:1
           RX packets:611 errors:0 dropped:0 overruns:0 frame:0
           TX packets:611 errors:0 dropped:0 overruns:0 carrier:0
           collisions:0 txqueuelen:1000
           RX bytes:66073 (64.5 KiB)  TX bytes:66073 (64.5 KiB)


```

17. Use the `exit` command to end the shell session and return to the terminal of the host machine.
18. To understand more fully how the shared networking model works in Docker, start an NGINX container in `host` network mode. The NGINX container automatically exposes port `80`, which we previously had to forward to a port on the host machine. Use the `docker run` command to start an NGINX container on the host machine:

```
docker run -itd --network host --name hostnet2 nginx:latest
```

This command will start an NGINX container in the `host` networking mode.

19. Navigate to `http://localhost:80` using a web browser on the host machine:



Welcome to nginx!

If you see this page, the nginx web server is successfully installed and working. Further configuration is required.

For online documentation and support please refer to nginx.org.
Commercial support is available at nginx.com.

Thank you for using nginx.

You should be able to see the NGINX default web page displayed in your web browser. It should be noted that the `docker run` command did not explicitly forward or expose any ports to the host machine. Since the container is running in `host` networking mode, any ports that containers expose by default will be available directly on the host machine.

20. Use the `docker run` command to create another NGINX instance in the `host` network mode. Call this container `hostnet3` to differentiate it from the other two container instances:

```
docker run -itd --network host --name hostnet3 nginx:latest
```

21. Now use the `docker ps -a` command to list all the containers, both in running and stopped status:

```
docker ps -a
```

The list of running containers will be displayed:

CONTAINER ID	IMAGE	COMMAND	CREATED
STATUS		PORTS	NAMES
da56fcf81d02	nginx:latest	"nginx -g 'daemon of...'"	4 minutes ago
Exited (1)	4 minutes ago		hostnet3
5786dac6fd27	nginx:latest	"nginx -g 'daemon of...'"	37 minutes ago
Up 37 minutes			hostnet2
648b291846e7	alpine:latest	"/bin/sh"	38 minutes ago
Up 38 minutes			hostnet

22. Based on the preceding output, you can see that the `hostnet3` container exited and is currently in a stopped state. To understand more fully why this is the case, use the `docker logs` command to view the container logs:

```
docker logs hostnet3
```

The log output should be displayed as follows:

```
[200~2020/05/15 00:50:04 [emerg] 1#1: bind() to 0.0.0.0:80 failed (98: Address already in use)
nginx: [emerg] bind() to 0.0.0.0:80 failed (98: Address already in use)
2020/05/15 00:50:04 [emerg] 1#1: bind() to 0.0.0.0:80 failed (98: Address already in use)
nginx: [emerg] bind() to 0.0.0.0:80 failed (98: Address already in use)
2020/05/15 00:50:04 [emerg] 1#1: bind() to 0.0.0.0:80 failed (98: Address already in use)
nginx: [emerg] bind() to 0.0.0.0:80 failed (98: Address already in use)
2020/05/15 00:50:04 [emerg] 1#1: bind() to 0.0.0.0:80 failed (98: Address already in use)
nginx: [emerg] bind() to 0.0.0.0:80 failed (98: Address already in use)
2020/05/15 00:50:04 [emerg] 1#1: bind() to 0.0.0.0:80 failed (98: Address already in use)
nginx: [emerg] bind() to 0.0.0.0:80 failed (98: Address already in use)
2020/05/15 00:50:04 [emerg] 1#1: still could not bind()
nginx: [emerg] still could not bind()
```

Essentially, this second instance of an NGINX container was unable to start properly because it was unable to bind to port `80` on the host machine. The reason for this is that the `hostnet2` container is already listening on that port.

37. Next, stop all running containers using `docker stop`, followed by the container name or ID:

```
docker stop hostnet1
```

Repeat this step for all running containers in your environment.

38. Clean up the container images and unused networks using the `docker system prune` command:

```
docker system prune -fa
```

This command will clean up all unused container images, networks, and volumes remaining on your machine.

In this exercise, we looked at the default networking drivers available by default in Docker: `bridge`, `host` and `none`. For each example, we explored how the network functions, how containers deployed using these network drivers function with the host machine, and how they function with other containers on the network.

Summary

In this lab, we looked at the many facets of networking in relation to microservices and Docker containers. Docker comes equipped with numerous drivers and configuration options that users can use to tune the way their container networking works in almost any environment.

In the next lab, we will look at the next pillar of a powerful containerized infrastructure: storage.